

CS 441 - Final Project

April 25, 2024

1 CS 441 Final Project - Disaster Tweets

1.1 1. Import

```
[ ]: import pandas as pd
import numpy as np
import nltk
from nltk.corpus import stopwords
from nltk.stem import WordNetLemmatizer
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, f1_score
import matplotlib.pyplot as plt
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
```

```
[ ]: train_df = pd.read_csv("train.csv")

# Split the dataset (80% train, 20% validation)
train_data, val_data = train_test_split(train_df, test_size = 0.2, random_state=
    ↳ 441)

print("Training set:", train_data.shape)
print("Validation set:", val_data.shape)

print("Example:\n", train_df.iloc[32])
```

Training set: (6090, 5)

Validation set: (1523, 5)

Example:

id	49
keyword	ablaze
location	Est. September 2012 - Bristol
text	We always try to bring the heavy. #metal #RT h...
target	0

Name: 32, dtype: object

2. Preprocess

```
[ ]: # Preprocess text, remove part of speech, stopwords, and http links
# https://stackoverflow.com/questions/17390326/
↳getting-rid-of-stop-words-and-document-tokenization-using-nltk
lemmatizer = WordNetLemmatizer()
stop_words = set(stopwords.words("english"))

def preprocess_text(text):
    tokens = nltk.word_tokenize(text.lower())
    tokens = [lemmatizer.lemmatize(token) for token in tokens if token.
↳isalpha() and token not in stop_words]
    return " ".join(tokens)

train_data["clean_text"] = train_data["text"].apply(preprocess_text)
val_data["clean_text"] = val_data["text"].apply(preprocess_text)

tfidf_vectorizer = TfidfVectorizer(max_features = 5000)

X_train = tfidf_vectorizer.fit_transform(train_data["clean_text"])
y_train = train_data["target"]

X_val = tfidf_vectorizer.transform(val_data["clean_text"])
y_val = val_data["target"]

print("Original:", train_data.iloc[1]["text"])
print("Processed:", train_data.iloc[1]["clean_text"])
```

Original: British diver Neil Anthony Fears found dead by the wreck of a steamship - Daily Mail <http://t.co/QP3GVvfoFq>
Processed: british diver neil anthony fear found dead wreck steamship daily mail http

2.1 3. KNN

```
[ ]: k_values = [1, 3, 5, 7, 9, 11]

accuracy_scores = []
f1_scores = []

for k in k_values:
    knn_model = KNeighborsClassifier(n_neighbors = k)
    knn_model.fit(X_train, y_train)

    y_pred = knn_model.predict(X_val)
    accuracy = accuracy_score(y_val, y_pred)
    f1 = f1_score(y_val, y_pred)
```

```

    accuracy_scores.append(accuracy)
    f1_scores.append(f1)

result_df = pd.DataFrame({
    'K': [K for K in k_values],
    'Accuracy': accuracy_scores,
    'F1-score': f1_scores
})
print(result_df)


```

	K	Accuracy	F1-score
0	1	0.695995	0.532795
1	3	0.670387	0.439732
2	5	0.647406	0.356886
3	7	0.640184	0.330073
4	9	0.633618	0.304239
5	11	0.625739	0.278481

```

[ ]: plt.plot(k_values, accuracy_scores, label='Accuracy')
plt.plot(k_values, f1_scores, label='F1-score')
plt.xlabel('Neighbors')
plt.ylabel('Score')
plt.title('KNN')
plt.xticks(k_values)
plt.legend()
plt.show()

```



2.2 4. Logistic Regression

```
[ ]: C_values = [0.01, 0.1, 1, 10, 100]
penalty_values = ['l1', 'l2']

hyperparam_values = []
accuracy_scores = []
f1_scores = []

for C in C_values:
    for penalty in penalty_values:
        lr_clf = LogisticRegression(C = C, penalty = penalty, solver = 'liblinear', random_state = 42)
        lr_clf.fit(X_train, y_train)

        y_pred = lr_clf.predict(X_val)
        accuracy = accuracy_score(y_val, y_pred)
        f1 = f1_score(y_val, y_pred)
```

```

        hyperparam_values.append((C, penalty))
        accuracy_scores.append(accuracy)
        f1_scores.append(f1)

results_df = pd.DataFrame({
    'C': [param[0] for param in hyperparam_values],
    'Penalty': [param[1] for param in hyperparam_values],
    'Accuracy': accuracy_scores,
    'F1-score': f1_scores
})

print(results_df)

```

	C	Penalty	Accuracy	F1-score
0	0.01	11	0.556796	0.000000
1	0.01	12	0.556796	0.000000
2	0.10	11	0.582403	0.260465
3	0.10	12	0.730138	0.587763
4	1.00	11	0.766907	0.705394
5	1.00	12	0.800394	0.755233
6	10.00	11	0.762968	0.721236
7	10.00	12	0.776756	0.736434
8	100.00	11	0.746553	0.706240
9	100.00	12	0.757715	0.718964

```

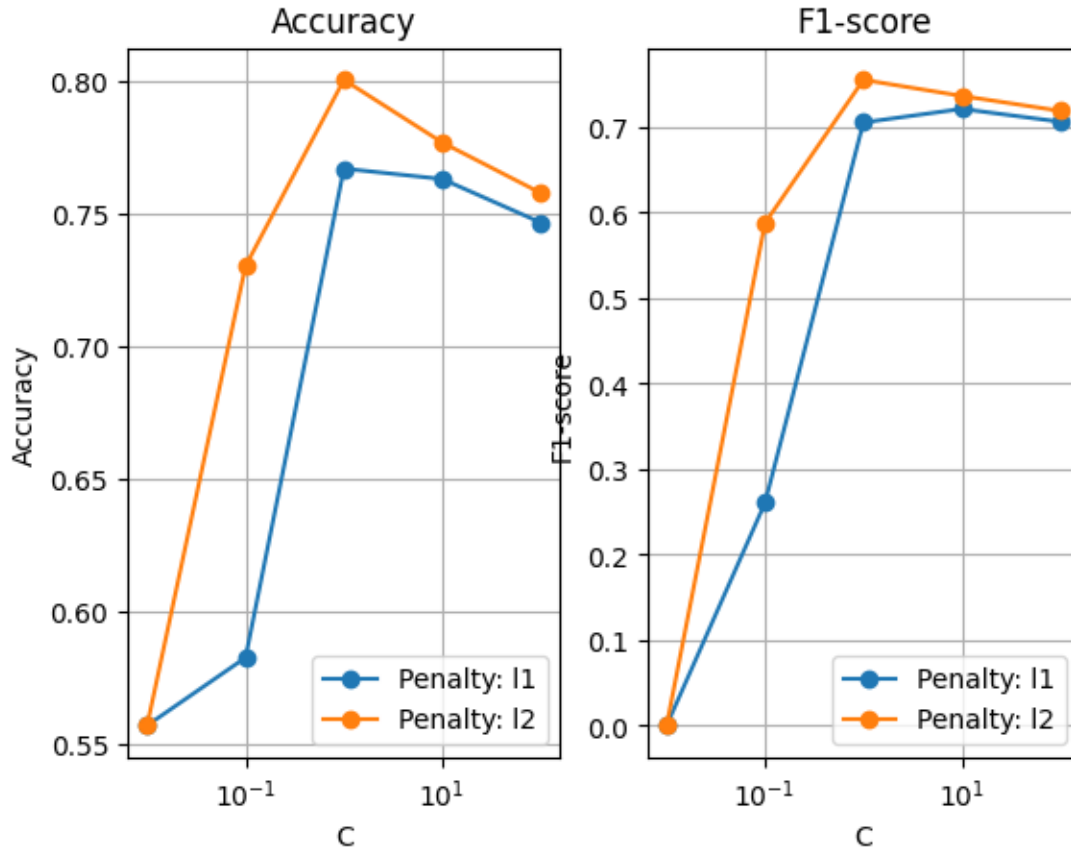
[ ]: plt.subplot(1, 2, 1)
    for penalty in penalty_values:
        penalty_results = results_df[results_df['Penalty'] == penalty]
        plt.plot(penalty_results['C'], penalty_results['Accuracy'], marker='o',
        label=f'Penalty: {penalty}')
    plt.xlabel('C')
    plt.ylabel('Accuracy')
    plt.title('Accuracy')
    plt.xscale('log')
    plt.legend()
    plt.grid(True)

plt.subplot(1, 2, 2)
    for penalty in penalty_values:
        penalty_results = results_df[results_df['Penalty'] == penalty]
        plt.plot(penalty_results['C'], penalty_results['F1-score'], marker='o',
        label = f'Penalty: {penalty}')
    plt.xlabel('C')
    plt.ylabel('F1-score')
    plt.title('F1-score')
    plt.xscale('log')

```

```
plt.legend()
plt.grid(True)

plt.show()
```



2.3 5. SVM

```
[ ]: C_values = [0.1, 1, 10, 100]
kernel_values = ['linear', 'rbf', 'poly']

hyperparam_values = []
accuracy_scores = []
f1_scores = []

for C in C_values:
    for kernel in kernel_values:
        svm_model = SVC(C = C, kernel = kernel, random_state = 441)
        svm_model.fit(X_train, y_train)
```

```

y_pred = svm_model.predict(X_val)
accuracy = accuracy_score(y_val, y_pred)
f1 = f1_score(y_val, y_pred)

hyperparam_values.append((C, kernel))
accuracy_scores.append(accuracy)
f1_scores.append(f1)

results_df = pd.DataFrame({
    'C': [param[0] for param in hyperparam_values],
    'Kernel': [param[1] for param in hyperparam_values],
    'Accuracy': accuracy_scores,
    'F1-score': f1_scores
})

print(results_df)

```

	C	Kernel	Accuracy	F1-score
0	0.1	linear	0.722259	0.561658
1	0.1	rbf	0.599475	0.177898
2	0.1	poly	0.587656	0.130194
3	1.0	linear	0.791858	0.749010
4	1.0	rbf	0.803020	0.748744
5	1.0	poly	0.717663	0.557613
6	10.0	linear	0.746553	0.706687
7	10.0	rbf	0.793828	0.747994
8	10.0	poly	0.727511	0.587885
9	100.0	linear	0.735391	0.693069
10	100.0	rbf	0.791858	0.744972
11	100.0	poly	0.711753	0.546956

```

[ ]: plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
for kernel in kernel_values:
    kernel_results = results_df[results_df['Kernel'] == kernel]
    plt.plot(kernel_results['C'], kernel_results['Accuracy'], marker='o',
             label=f'Kernel: {kernel}')
plt.xlabel('C')
plt.ylabel('Accuracy')
plt.title('Accuracy')
plt.xscale('log')
plt.legend()
plt.grid(True)

plt.subplot(1, 2, 2)
for kernel in kernel_values:
    kernel_results = results_df[results_df['Kernel'] == kernel]


```

```

plt.plot(kernel_results['C'], kernel_results['F1-score'], marker='o',
         label=f'Kernel: {kernel}')
plt.xlabel('C')
plt.ylabel('F1-score')
plt.title('F1-score')
plt.xscale('log')
plt.legend()
plt.grid(True)

plt.show()

```



3 6. Test

```

[ ]: train_df = pd.read_csv("train.csv")
train_df["clean_text"] = train_df["text"].apply(preprocess_text)
X_train = tfidf_vectorizer.transform(train_df["clean_text"])
y_train = train_df["target"]

test_df = pd.read_csv("test.csv")
test_df["clean_text"] = test_df["text"].apply(preprocess_text)
X_test = tfidf_vectorizer.transform(test_df["clean_text"])

svm_model = SVC(C = 1.0, kernel = 'linear', random_state = 441)
svm_model.fit(X_train, y_train)
y_test = svm_model.predict(X_test)

```



```
output = pd.DataFrame({'id': test_df['id']})
output['target'] = y_test
```

```
[ ]: print(output)
      output.to_csv("result.csv", index = False)
```

	id	target
0	0	1
1	2	1
2	3	1
3	9	1
4	11	1
...
3258	10861	1
3259	10865	1
3260	10868	1
3261	10874	1
3262	10875	0

[3263 rows x 2 columns]